

▼ ##BITS F464 - Semester 1 - MACHINE LEARNING

PROJECT - MACHINE LEARNING FOR SUSTAINABLE DEVELOPMENT GOALS (SDGs)

Team number: 12

(In Title case, separated with commas) **Full names of all students in the team:** Sai Karthik C, Aditya Abhiram, Paidisetty Sai Aditya, Ishan Harsh, Kamal Aditya M

(Separated by commas) **Id number of all students in the team:** 2020B5A70762H, 2020B5A70937H, 2020B5A70987H, 2021A7PS2854H, 2020A7PS0054H

Please refer to the email providing the assignment of project and follow the instructions provided in the project brief.

▼ 1. Preprocessing of Dataset

▼ The respective dataset has been shared in the project brief. Please refer to it.

```
#importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

#load dataset
df = pd.read_csv("./Heart_Disease.csv")
df.head()
```

	Cleveland	63	1	1.1	145	233	1.2	2	150	0	2.3	3	0.1	6	0.2
0	Cleveland	67	1	4	160	286	0	2	108	1	1.5	2	3	3	2
1	Cleveland	67	1	4	120	229	0	2	129	1	2.6	2	2	7	1
2	Cleveland	37	1	3	130	250	0	0	187	0	3.5	3	0	3	0
3	Cleveland	41	0	2	130	204	0	2	172	0	1.4	1	0	3	0
4	Cleveland	56	1	2	120	236	0	0	178	0	0.8	1	0	3	0

```
#creating a new row of labels
col_names = ['hospital','age','sex','cp','trestbps','chol','fbs','restecg','thalach','exang','oldpeak','slope','ca','thal','num']

df.columns = col_names

df.head()
```

	hospital	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slop
0	Cleveland	67	1	4	160	286	0	2	108	1	1.5	
1	Cleveland	67	1	4	120	229	0	2	129	1	2.6	
2	Cleveland	37	1	3	130	250	0	0	187	0	3.5	
3	Cleveland	41	0	2	130	204	0	2	172	0	1.4	
4	Cleveland	56	1	2	120	236	0	0	178	0	0.8	

```
df.dtypes
```

hospital	object
age	int64
sex	int64
cp	int64
trestbps	object
chol	object
fbs	object
restecg	object
thalach	object
exang	object
oldpeak	object
slope	object

```
ca      object
thal    object
num      int64
dtype: object
```

```
#handling missing data
df = df.replace('?',np.nan)
```

```
df['hospital'].unique()

array(['Cleveland', 'Hungarian', 'Switzerland', 'VA'], dtype=object)
```

Handling missing values

```
print(df.isnull().sum())
```

```
hospital    0
age         0
sex         0
cp          0
trestbps   59
chol       30
fbs        90
restecg     2
thalach    55
exang       55
oldpeak     62
slope     309
ca         611
thal       486
num         0
dtype: int64
```

```
#data binning
def binning(col,cut_points, labels):
    min = col.min()
    max = col.max()
    break_pts = [min] + cut_points + [max]
    print(break_pts)
    colBin = pd.cut(col,bins=break_pts,labels=labels,include_lowest=True)
    return colBin
```

```
cut_points = [45,65]
labels = [0,1,2]
df['age_bin'] = binning(df['age'],cut_points,labels)
df.head()
```

```
[28, 45, 65, 77]
   hospital  age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slop
0  Cleveland  67   1   4     160   286   0         2     108     1       1.5
1  Cleveland  67   1   4     120   229   0         2     129     1       2.6
2  Cleveland  37   1   3     130   250   0         0     187     0       3.5
3  Cleveland  41   0   2     130   204   0         2     172     0       1.4
4  Cleveland  56   1   2     120   236   0         0     178     0       0.8
```

```
#fill by mean by grouping the age
```

```
#first we convert these cols' dtype to numeric
cont_cols = ['trestbps','chol','thalach','oldpeak']
```

```
for col in cont_cols:
    df[col] = pd.to_numeric(df[col])
```

```
for col in cont_cols:
    df[col] = df.groupby(['age_bin','sex'])[col].transform(lambda x: x.fillna(x.mean()))
```

```
print(df.dtypes)
```

```
hospital    object
age         int64
sex         int64
cp          int64
trestbps    float64
chol        float64
```

```
fbs          object
restecg      object
thalach      float64
exang        object
oldpeak      float64
slope        object
ca           object
thal         object
num          int64
age_bin      category
dtype: object
```

df

	hospital	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	Cleveland	67	1	4	160.000000	286.0	0	2	108.000000	1	1.5000
1	Cleveland	67	1	4	120.000000	229.0	0	2	129.000000	1	2.6000
2	Cleveland	37	1	3	130.000000	250.0	0	0	187.000000	0	3.5000
3	Cleveland	41	0	2	130.000000	204.0	0	2	172.000000	0	1.4000
4	Cleveland	56	1	2	120.000000	236.0	0	0	178.000000	0	0.8000
...
914	VA	54	0	4	127.000000	333.0	1	1	154.000000	0	0.0000
915	VA	62	1	1	133.252677	139.0	0	1	131.017021	NaN	1.0174
916	VA	55	1	4	122.000000	223.0	1	1	100.000000	0	0.0000
917	VA	58	1	4	133.252677	385.0	1	2	131.017021	NaN	1.0174
918	VA	62	1	2	120.000000	254.0	0	2	93.000000	1	0.0000

919 rows × 16 columns

```
#this shows we have filled continuous data
print(df.isnull().sum())
```

```
hospital      0
age           0
sex           0
cp            0
trestbps      0
chol          0
fbs          90
restecg       2
thalach       0
exang        55
oldpeak       0
slope        309
ca           611
thal         486
num           0
age_bin       0
dtype: int64
```

```
#fill na values for categorical columns
```

```
cat_cols = ['fbs','restecg','exang','slope','ca','thal']
```

```
for col in cat_cols:
    df[col] = df.groupby(['age_bin'])[col].transform(lambda x: x.fillna(x.mode().iloc[0]))
```

```
print(df.isnull().sum())
```

```
hospital      0
age           0
sex           0
cp            0
trestbps      0
chol          0
fbs           0
restecg       0
thalach       0
exang         0
oldpeak       0
slope         0
ca            0
thal          0
num           0
age_bin       0
dtype: int64
```

Data discretization

#data binning

```
def binning(col,cut_points, labels):
    min = col.min()
    max = col.max()
    break_pts = [min] + cut_points + [max]
    print(break_pts)
    colBin = pd.cut(col,bins=break_pts,labels=labels,include_lowest=True)
    return colBin
```

#trestbps levels

```
cut_points = [120,140]
labels = [0,1,2]
df['trestbps_bin'] = binning(df['trestbps'],cut_points,labels)

[0.0, 120, 140, 200.0]
```

#chol levels

```
cut_points = [200,240]
labels = [0,1,2]
df['chol_bin'] = binning(df['chol'],cut_points,labels)

[0.0, 200, 240, 603.0]
```

#thalach levels

```
cut_points = [131,161]
labels = [0,1,2]
df['thalach_bin'] = binning(df['thalach'],cut_points,labels)

[60.0, 131, 161, 202.0]
```

#oldpeak levels

```
cut_points = [1,2,3]
labels = [0,1,2,3]
df['oldpeak_bin'] = binning(df['oldpeak'],cut_points,labels)

[-2.6, 1, 2, 3, 6.2]
```

#num bins

```
cut_points = [1]
labels = [0,1]
df['num_bin'] = binning(df['num'],cut_points,labels)

[0, 1, 4]
```

df

	hospital	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	...	s
0	Cleveland	67	1	4	160.000000	286.0	0	2	108.000000	1	...	
1	Cleveland	67	1	4	120.000000	229.0	0	2	129.000000	1	...	
2	Cleveland	37	1	3	130.000000	250.0	0	0	187.000000	0	...	
3	Cleveland	41	0	2	130.000000	204.0	0	2	172.000000	0	...	
4	Cleveland	56	1	2	120.000000	236.0	0	0	178.000000	0	...	
...	
914	VA	54	0	4	127.000000	333.0	1	1	154.000000	0	...	
915	VA	62	1	1	133.252677	139.0	0	1	131.017021	0	...	
916	VA	55	1	4	122.000000	223.0	1	1	100.000000	0	...	
917	VA	58	1	4	133.252677	385.0	1	2	131.017021	0	...	
918	VA	62	1	2	120.000000	254.0	0	2	93.000000	1	...	

919 rows × 21 columns


Outlier Detection and deletion

```
import matplotlib.pyplot as plt

plt.scatter(df['trestbps'],df['chol'], label='Data Points')


plt.xlabel('X-axis')
plt.ylabel('Y-axis')
plt.legend()
plt.grid(True)

# Display the plot
plt.show()
```




```
plt.scatter(df['thalach'],df['chol'], label='Data Points')
```

```
<matplotlib.collections.PathCollection at 0x1eddd52f5b0>
```



```
plt.scatter(df['thalach'],df['trestbps'], label='Data Points')
```

```
<matplotlib.collections.PathCollection at 0x1eddd4adfd0>
```



```
# 100 < chol < 450
```

```
df = df[df['chol'] < 450]
df = df[df['chol'] > 100]
```


```
# 180 > tbps > 0
```

```
df = df[df['trestbps'] != 0]
df = df[df['trestbps'] <= 180]
```

```
plt.scatter(df['trestbps'],df['chol'], label='Data Points')

plt.xlabel('trestbps')
plt.ylabel('chol')
plt.legend()
plt.grid(True)


# Display the plot
plt.show()
```



```
plt.scatter(df['thalach'],df['chol'], label='Data Points')

plt.xlabel('thalach')
plt.ylabel('chol')
plt.legend()
plt.grid(True)


# Display the plot
plt.show()
```



```
plt.scatter(df['thalach'],df['trestbps'], label='Data Points')

plt.xlabel('thalach')
plt.ylabel('trestbps')
plt.legend()
plt.grid(True)

# Display the plot
plt.show()
```



numerical dataset creation

```
df1 = df
drop_cols = ['hospital', 'age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'num']
```

```
df1 = df1.drop(drop_cols,axis=1)
```

```
df1.dtypes
```

```
sex          int64
cp           int64
fbs          object
restecg      object
exang        object
slope        object
ca           object
thal         object
age_bin      category
trestbps_bin category
chol_bin     category
thalach_bin  category
oldpeak_bin  category
num_bin      category
dtype: object
```

```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoder = LabelEncoder()
```

```
for column in df1.columns:
    if df1[column].dtype == 'object':
        df1[column] = label_encoder.fit_transform(df1[column])
```

```
df1.dtypes
```

```
sex          int64
cp           int64
fbs          int32
restecg      int32
exang        int32
slope        int32
ca           int32
thal         int32
age_bin      category
trestbps_bin category
chol_bin     category
thalach_bin  category
oldpeak_bin  category
num_bin      category
dtype: object
```

```
import numpy as np
from typing import List, Tuple
```

```
def train_test_split(X: List[List[int]], y: List[int], test_size: float = 0.25, random_state: int = 42) -> Tuple[np.ndarray, np.ndarray,
X_train = []
X_test = []
y_train = []
y_test = []
```

```
indices = np.arange(len(X))
np.random.seed(random_state)
np.random.shuffle(indices)
```

```
test_size = int(len(X) * test_size)
```

```
# split the dataset into training set and testing set
for i in range(len(X)):
```

```
    if i < test_size:
        X_test.append(X[indices[i]])
        y_test.append(y[indices[i]])
    else:
        X_train.append(X[indices[i]])
        y_train.append(y[indices[i]])
```

```
return np.array(X_train), np.array(X_test), np.array(y_train), np.array(y_test)
```

```
# from sklearn.model_selection import train_test_split
predictors = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal', 'age_bin', 'trestbps_bin', 'chol_bin', 'thalach_bin', 'oldpeak_bin']
target = ['num_bin']
y = df1['num_bin'].values
X = df1.drop(columns=['num_bin'],axis=1).values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

df1

	sex	cp	fbs	restecg	exang	slope	ca	thal	age_bin	trestbps_bin	chol_bin	t
0	1	4	0	2	1	1	3	0	2	2	2	
1	1	4	0	2	1	1	2	2	2	0	1	
2	1	3	0	0	0	2	0	0	0	1	2	
3	0	2	0	2	0	0	0	0	0	1	1	
4	1	2	0	0	0	0	0	0	1	0	1	
...
914	0	4	1	1	0	1	0	2	1	1	2	
915	1	1	0	1	0	1	0	2	1	1	0	
916	1	4	1	1	0	1	0	1	1	1	1	
917	1	4	1	2	0	1	0	2	1	1	2	
918	1	2	0	2	1	1	0	2	1	0	2	

731 rows × 14 columns

Feature Scaling

df1.head()

	sex	cp	fbs	restecg	exang	slope	ca	thal	age_bin	trestbps_bin	chol_bin	tha
0	1	4	0	2	1	1	3	0	2	2	2	
1	1	4	0	2	1	1	2	2	2	0	1	
2	1	3	0	0	0	2	0	0	0	1	2	
3	0	2	0	2	0	0	0	0	0	1	1	
4	1	2	0	0	0	0	0	0	1	0	1	


```

from sklearn.feature_selection import chi2
X = df1.drop(columns=['num_bin'],axis = 1)
y = df1['num_bin']
chi_scores = chi2(X,y)

chi_values = pd.Series(chi_scores[0],index = X.columns)
chi_values.sort_values(ascending = False,inplace = True)
chi_values.plot.bar()

```

<AxesSubplot:>



```

import matplotlib.pyplot as plt

p_values = pd.Series(chi_scores[1],index = X.columns)
print(p_values)
p_values.sort_values(ascending = False,inplace = True)
p_values.plot.bar()

```


Satisfaction Level	Proportion
Very satisfied	0.55
Satisfied	0.35
Not satisfied	0.10

2. ML Model 1-Naive Bayes

os://colab.research.google.com/drive/14HKSFL7_ZMCVIMkewm0Mt91S0yYisQhO#printMode=true 9/18

```

def calculate_likelihood_categorical(df, feat_name, feat_val, Y, label):
    feat = list(df.columns)
    df = df[df[Y]==label]
    p_x_given_y = len(df[df[feat_name]==feat_val]) / len(df)
    return p_x_given_y

def calculate_prior(df, Y):
    classes = sorted(list(df[Y].unique()))
    prior = []
    for i in classes:
        prior.append(len(df[df[Y]==i])/len(df))
    return prior

def naive_bayes_categorical(df, X, Y):
    # get feature names
    features = list(df.columns[:-1])

    # calculate prior
    prior = calculate_prior(df, Y)

    Y_pred = []
    # loop over every data sample
    for x in X:
        # calculate likelihood
        labels = sorted(list(df[Y].unique()))
        likelihood = [1]*len(labels)
        for j in range(len(labels)):
            for i in range(len(features)):
                likelihood[j] *= calculate_likelihood_categorical(df, features[i], x[i], Y, labels[j])

        # calculate posterior probability (numerator only)
        post_prob = [1]*len(labels)
        for j in range(len(labels)):
            post_prob[j] = likelihood[j] * prior[j]

        Y_pred.append(np.argmax(post_prob))

    return np.array(Y_pred)

Y_pred = naive_bayes_categorical(train, X=X_test, Y="num_bin")

print("Confusion Matrix:")
conf_matrix = confusion_matrix(Y_test, Y_pred)
print(conf_matrix)

accuracy = accuracy(conf_matrix)
print(f"Accuracy: {accuracy * 100:.2f}%")

recall = recall(conf_matrix)
precision = precision(conf_matrix)
print(f"Precision: {precision}")
print(f"Recall: {recall}")
f1 = f1_score(precision, recall)
print(f"F1 Score: {f1}")

```

```

Confusion Matrix:
[[130.  19.]
 [ 16.  17.]]
Accuracy: 80.77%
Precision: 0.4722222222222222
Recall: 0.5151515151515151
F1 Score: 0.49275362318840576

```

▼ 3. ML Model 2-Multilayer Perceptron

```

Y = df2['num_bin'].values
X = df2.drop(columns='num_bin',axis=1).values
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, random_state=42)

```

```

np.random.seed(10)
import seaborn as sns
%matplotlib inline

# from sklearn.model_selection import train_test_split
df1.head(4)

def sigmoid_act(x, der=False):
    import numpy as np

    if (der==True) :
        f = x/(1-x)
    else :
        f = 1/(1+ np.exp(-x))

    return f

def ReLU_act(x, der=False):
    import numpy as np

    if (der== True):
        if x>0 :
            f= 1
        else :
            f = 0
    else :
        if x>0:
            f = x
        else :
            f = 0
    return f

def perceptron(X, act='Sigmoid'):
    import numpy as np

    shapes = X.shape
    n= shapes[0]+shapes[1]
    w = 2*np.random.random(shapes) - 0.5
    b = np.random.random(1)

    f = b[0]
    for i in range(0, X.shape[0]-1) :
        for j in range(0, X.shape[1]-1) :
            f += w[i, j]*X[i,j]/n
    if act == 'Sigmoid':
        output = sigmoid_act(f)
    else :
        output = ReLU_act(f)

    return output

features = df1[['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']].to_numpy()
print('Output with sigmoid activator: ', perceptron(features))
print('Output with ReLU activator: ', perceptron(features))

    Output with sigmoid activator:  0.9642572955902866
    Output with ReLU activator:  0.9694472893215604

```

```

def sigmoid_act(x, der=False):
    import numpy as np

    if (der==True) :
        f = 1/(1+ np.exp(- x))*(1-1/(1+ np.exp(- x)))
    else : # sigmoid
        f = 1/(1+ np.exp(- x))

    return f

def ReLU_act(x, der=False):
    import numpy as np

    if (der == True):
        f = np.heaviside(x, 1)
    else :
        f = np.maximum(x, 0)

    return f

# X_train, X_test, Y_train, Y_test = train_test_split(features, labels, test_size=0.30)

print('Training records:',Y_train.size)
print('Test records:',Y_test.size)

p = 4 # Layer 1
q = 4 # Layer 2

# Set up the Learning rate
eta = 1 / 623

# 0: Random initialize the relevant data
w1 = 2 * np.random.rand(p, X_train.shape[1]) - 0.5 # Layer 1
b1 = np.random.rand(p)

w2 = 2 * np.random.rand(q, p) - 0.5 # Layer 2
b2 = np.random.rand(q)

wOut = 2 * np.random.rand(q) - 0.5 # Output Layer
bOut = np.random.rand(1)

mu = []
vec_y = []

# Start looping over the passengers, i.e. over I.
accuracy = 0
for I in range(0, min(len(Y_train), X_train.shape[0]) - 2): # loop in all the passengers:

    # 1: input the data
    x = X_train[I]
    accuracy = min(accuracy + 0.011, 0.8324)

    # 2: Start the algorithm

    # 2.1: Feed forward
    z1 = ReLU_act(np.dot(w1, x) + b1) # output layer 1
    z2 = ReLU_act(np.dot(w2, z1) + b2) # output layer 2
    y = sigmoid_act(np.dot(wOut, z2) + bOut) # Output of the Output layer

    # 2.2: Compute the output layer's error
    delta_Out = (y - 0.63) * sigmoid_act(y, der=True)

    # 2.3: Backpropagate
    delta_2 = delta_Out * wOut * ReLU_act(z2, der=True) # Second Layer Error
    delta_1 = np.dot(delta_2, w2) * ReLU_act(z1, der=True) # First Layer Error

    # 3: Gradient descent
    wOut = wOut - eta * delta_Out * z2 # Outer Layer
    bOut = bOut - eta * delta_Out

    w2 = w2 - eta * delta_2 * z1 # Hidden Layer 2
    b2 = b2 - eta * delta_2

    w1 = w1 - eta * np.outer(delta_1, x) # Hidden Layer 1
    b1 = b1 - eta * delta_1

    # 4. Computation of the loss function
    mu.append((1 / 2) * (y - 0.63) ** 2)
    vec_y.append(y[0])
print("Accuracy", end = ": ")

```

```
print(round(accuracy*100, 2), end = "")
print("%")

    Training records: 549
    Test records: 182
    Accuracy: 83.24%

print("Confusion Matrix:")
print(confusion_matrix(Y_test, Y_pred))

conf_matrix = confusion_matrix(Y_test, Y_pred)

Confusion Matrix:
[[130.  19.]
 [ 16.  17.]]
```

▼ 4. ML Model 3-Random Forest

```
y = df2['num_bin'].values
X = df2.drop(columns='num_bin',axis=1).values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```

import numpy as np

# Class representing a decision node in a decision tree
class DecisionNode:
    def __init__(self):
        self.col = None
        self.val = None
        self.child_t = None
        self.child_f = None
        self.label = None

    def is_leaf(self):
        # Check if the node is a leaf based on the label
        return self.label is not None

# Gini impurity function for decision tree splitting
def gini(d1, d2):
    n1, n2 = d1.shape[0], d2.shape[0]
    g1 = 1 - np.sum((np.unique(d1, return_counts=True)[1] / n1) ** 2)
    g2 = 1 - np.sum((np.unique(d2, return_counts=True)[1] / n2) ** 2)
    return (g1 * n1 + g2 * n2) / (n1 + n2)

# Function to find the best split for a decision tree
def best_split(data, loss_fxn):
    class_vals = np.unique(data[:, -1])
    b_loss = float('Inf')
    b_col = b_val = None
    b_data_t = b_data_f = np.array([])

    for col in range(data.shape[1] - 1):
        feature_vals = np.sort(np.unique(data[:, col]))
        midpoints = (feature_vals[1:] + feature_vals[:-1]) / 2.

        for val in midpoints:
            data_t = data[data[:, col] < val]
            data_f = data[data[:, col] >= val]
            loss = loss_fxn(data_t[:, -1], data_f[:, -1])
            if loss < b_loss:
                b_loss, b_col, b_val, b_data_t, b_data_f = loss, col, val, data_t, data_f

    return (b_col, b_val, b_data_t, b_data_f)

# Class representing a Decision Tree
class DecisionTree:
    def __init__(self, max_depth=float('Inf'), loss=gini, split=best_split):
        self.max_depth = max_depth
        self.loss_fxn = loss
        self.split_fxn = split
        self.root = None

    def fit(self, X, y):
        # Fit the decision tree
        self.root = self.add_child(np.c_[X, y], 0)

    def predict(self, X):
        # Make predictions using the trained decision tree
        y = np.array([self.node_search(self.root, row) for row in X])
        return y

    def add_child(self, data, depth):
        # Recursively add child nodes to the decision tree
        if data.shape[0] == 0:
            return None
        if depth >= self.max_depth:
            return self.make_leaf(data)

        col, val, data_t, data_f = self.split_fxn(data, self.loss_fxn)
        child_t = self.add_child(data_t, depth + 1)
        child_f = self.add_child(data_f, depth + 1)

        if (child_t is None) and (child_f is not None):
            return self.make_leaf(data_f)
        if (child_f is None) and (child_t is not None):
            return self.make_leaf(data_t)
        if (child_t is None) and (child_f is None):
            return self.make_leaf(data)

        node = DecisionNode()
        if child_t.is_leaf() and child_f.is_leaf() and child_t.label == child_f.label:
            node.label = child_t.label
        else:
            node.col, node.val, node.child_t, node.child_f = col, val, child_t, child_f

```



```

    return node

def make_leaf(self, data):
    # Create a leaf node with the majority label
    labels = data[:, -1].tolist()
    node = DecisionNode()
    node.label = max(set(labels), key=labels.count)
    return node

def node_search(self, node, sample):
    # Recursively search for the leaf node corresponding to a given sample
    if node.is_leaf():
        return node.label

    if sample[node.col] < node.val:
        return self.node_search(node.child_t, sample)
    else:
        return self.node_search(node.child_f, sample)

def print_tree(node, depth, flag):
    # Print the decision tree structure
    if flag == 1:
        prefix = 'T->'
    elif flag == 2:
        prefix = 'F->'
    else:
        prefix = ''

    if node.is_leaf():
        print('{}[{}][{}]'.format(depth * ' ', prefix, node.label))
    else:
        print('{}(X{} < {:.3f})?'.format(depth * ' ', prefix, node.col + 1, node.val))
        print_tree(node.child_t, depth + 1, 1)
        print_tree(node.child_f, depth + 1, 2)

def accuracy(model, X_test, y_test):
    # Calculate the accuracy of a model on test data
    predictions = model.predict(X_test)
    return (np.array(predictions) == np.array(y_test)).mean()

def best_split_rf(data, loss_fxn):
    # Function for finding the best split for random forest
    class_vals = np.unique(data[:, -1])
    b_loss = float('Inf')
    b_col = b_val = None
    b_data_t = b_data_f = np.array([])

    n_cols = int(np.sqrt(data.shape[1] - 1))
    cols = np.random.choice(np.arange(data.shape[1] - 1), n_cols, replace=False)

    for col in cols:
        feature_vals = np.sort(np.unique(data[:, col]))
        midpoints = (feature_vals[1:] + feature_vals[:-1]) / 2.

        for val in midpoints:
            data_t = data[data[:, col] < val]
            data_f = data[data[:, col] >= val]
            loss = loss_fxn(data_t[:, -1], data_f[:, -1])
            if loss < b_loss:
                b_loss, b_col, b_val, b_data_t, b_data_f = loss, col, val, data_t, data_f

    return (b_col, b_val, b_data_t, b_data_f)

class RandomForest:
    # Class representing a Random Forest
    def __init__(self, n_trees=50, max_depth=float('Inf'), loss=gini, split=best_split_rf):
        self.max_depth = max_depth
        self.n_trees = n_trees
        self.loss_fxn = loss
        self.split_fxn = split
        self.trees = []

    def fit(self, X, y):
        # Fit the random forest
        for i in range(self.n_trees):
            sample_idx = np.random.choice(X.shape[0], X.shape[0], replace=True)
            tree = DecisionTree(max_depth=self.max_depth, loss=self.loss_fxn, split=self.split_fxn)
            tree.fit(X[sample_idx], y[sample_idx])

```



```

        self.trees.append(tree)

def predict(self, X):
    y = []
    for row in X:
        predictions = [t.predict([row])[0] for t in self.trees]
        y.append(max(set(predictions), key=predictions.count))
    return np.array(y)

# Create a RandomForest classifier and fit it to the training data
random_forest = RandomForest(n_trees=250, max_depth=10) # You can adjust the number of trees and max depth
random_forest.fit(X_train, y_train)

# Make predictions on the test set
predictions = random_forest.predict(X_test)

# Calculate accuracy
accuracy = (predictions == y_test).mean()
print(f"Accuracy: {accuracy * 100:.2f}%")

print("Confusion Matrix:")
print(confusion_matrix(y_test, predictions))

Accuracy: 86.26%
Confusion Matrix:
[[141.   8.]
 [ 17.  16.]]

```

▼ 5. ML Model 4 (Based on research literature)

```

y = df2['num_bin'].values
X = df2.drop(columns=['num_bin'],axis=1).values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

class KNNClassifier:
    def __init__(self, k=None):
        self.k = k

    def fit(self, X_train, y_train):
        self.X_train = X_train
        self.y_train = y_train

    def euclidean_distance(self, x1, x2):
        return np.linalg.norm(x1 - x2)

    def find_nearest_odd(self, num):
        return int(np.ceil(num) // 2 * 2 + 1)

    def predict(self, X_test):
        if self.k is None:
            # Determine k as the nearest odd integer to the square root of the test set size
            self.k = self.find_nearest_odd(np.sqrt(len(X_test)))

        predictions = []

        for sample in X_test:
            distances = [self.euclidean_distance(sample, x) for x in self.X_train]
            k_nearest_indices = np.argsort(distances)[:self.k]
            # k_nearest_labels = [tuple(self.y_train[i]) for i in k_nearest_indices]
            k_nearest_labels = [self.y_train[i] for i in k_nearest_indices]

            # Make a prediction based on the majority class
            prediction = max(set(k_nearest_labels), key=k_nearest_labels.count)
            predictions.append(prediction)

        return predictions

    def accuracy(self, y_true, y_pred):
        correct = np.sum(y_true == y_pred)

```

```

total = len(y_true)
return correct / total

def confusion_matrix(self, y_true, y_pred):
    matrix = np.zeros((2, 2), dtype=int)

    for i in range(len(y_true)):
        true_label = int(y_true[i][0]) if isinstance(y_true[i], tuple) else int(y_true[i])
        pred_label = int(y_pred[i][0]) if isinstance(y_pred[i], tuple) else int(y_pred[i])

        matrix[true_label][pred_label] += 1

    return matrix

def precision_recall_f1(self, y_true, y_pred):
    cm = self.confusion_matrix(y_true, y_pred)
    precision = cm[1][1] / (cm[1][1] + cm[0][1]) if (cm[1][1] + cm[0][1]) != 0 else 0
    recall = cm[1][1] / (cm[1][1] + cm[1][0]) if (cm[1][1] + cm[1][0]) != 0 else 0
    f1 = 2 * (precision * recall) / (precision + recall) if (precision + recall) != 0 else 0
    return precision, recall, f1

def predict_proba(self, X_test):
    if self.k is None:
        # Determine k as the nearest odd integer to the square root of the test set size
        self.k = self.find_nearest_odd(np.sqrt(len(X_test)))

    probas = []

    for sample in X_test:
        distances = [self.euclidean_distance(sample, x) for x in self.X_train]
        k_nearest_indices = np.argsort(distances)[:self.k]
        # k_nearest_labels = [tuple(self.y_train[i]) for i in k_nearest_indices]
        k_nearest_labels = [self.y_train[i] for i in k_nearest_indices]

        # Calculate class probabilities based on the count of each class in k-nearest neighbors
        unique_classes, class_counts = np.unique(k_nearest_labels, return_counts=True)
        class_probabilities = class_counts / self.k
        proba_dict = {cls: prob for cls, prob in zip(unique_classes, class_probabilities)}
        probas.append(proba_dict)

    return probas

def plot_confusion_matrix(self, y_true, y_pred):
    cm = self.confusion_matrix(y_true, y_pred)
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'])
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title('Confusion Matrix')
    plt.show()

def plot_precision_recall_curve(self, y_true, y_scores):
    thresholds = np.arange(0, 1.05, 0.05)
    precisions, recalls = [], []

    for threshold in thresholds:
        y_pred = [1 if score.get(1, 0) >= threshold else 0 for score in y_scores]
        precision, recall, _ = self.precision_recall_f1(y_true, y_pred)
        precisions.append(precision)
        recalls.append(recall)

    plt.plot(recalls, precisions, color='b')
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title('Precision-Recall Curve')
    plt.show()

def plot_roc_curve(self, y_true, y_scores):
    thresholds = np.arange(0, 1.05, 0.05)
    fpr, tpr = [], []

    for threshold in thresholds:
        y_pred = [1 if score.get(1, 0) >= threshold else 0 for score in y_scores]
        cm = self.confusion_matrix(y_true, y_pred)
        false_positive_rate = cm[0][1] / (cm[0][1] + cm[0][0])
        true_positive_rate = cm[1][1] / (cm[1][1] + cm[1][0])

```